

Solution Architecture for Gen AI



Ram N Sangwan

- Solution Guidelines - Well-Architected principles
- Chat Session management
- Architectures for Various Use Cases
- Manage Token limitation
- Deployment Standards – Cloud or On-Prem
- Local Chat

Solution Guidelines - Well-Architected principles

Typically focus on key areas to ensure the technology is used effectively and responsibly.

Some insights:

Choice of Foundation Model: depends domain, data, and the specific use case.

Accessibility and Control: on Prem for full control or using them as a managed cloud service for speed and simplicity.

The decision will impact factors like infrastructure management, cost predictability, and complexity.

Solution Guidelines - Well-Architected principles

Architecture Components: data processing layer, a generative model layer, a feedback & improvement layer, and a deployment & integration layer.

Operational and Model Risks: Models must be managed to control privacy and ensure security.

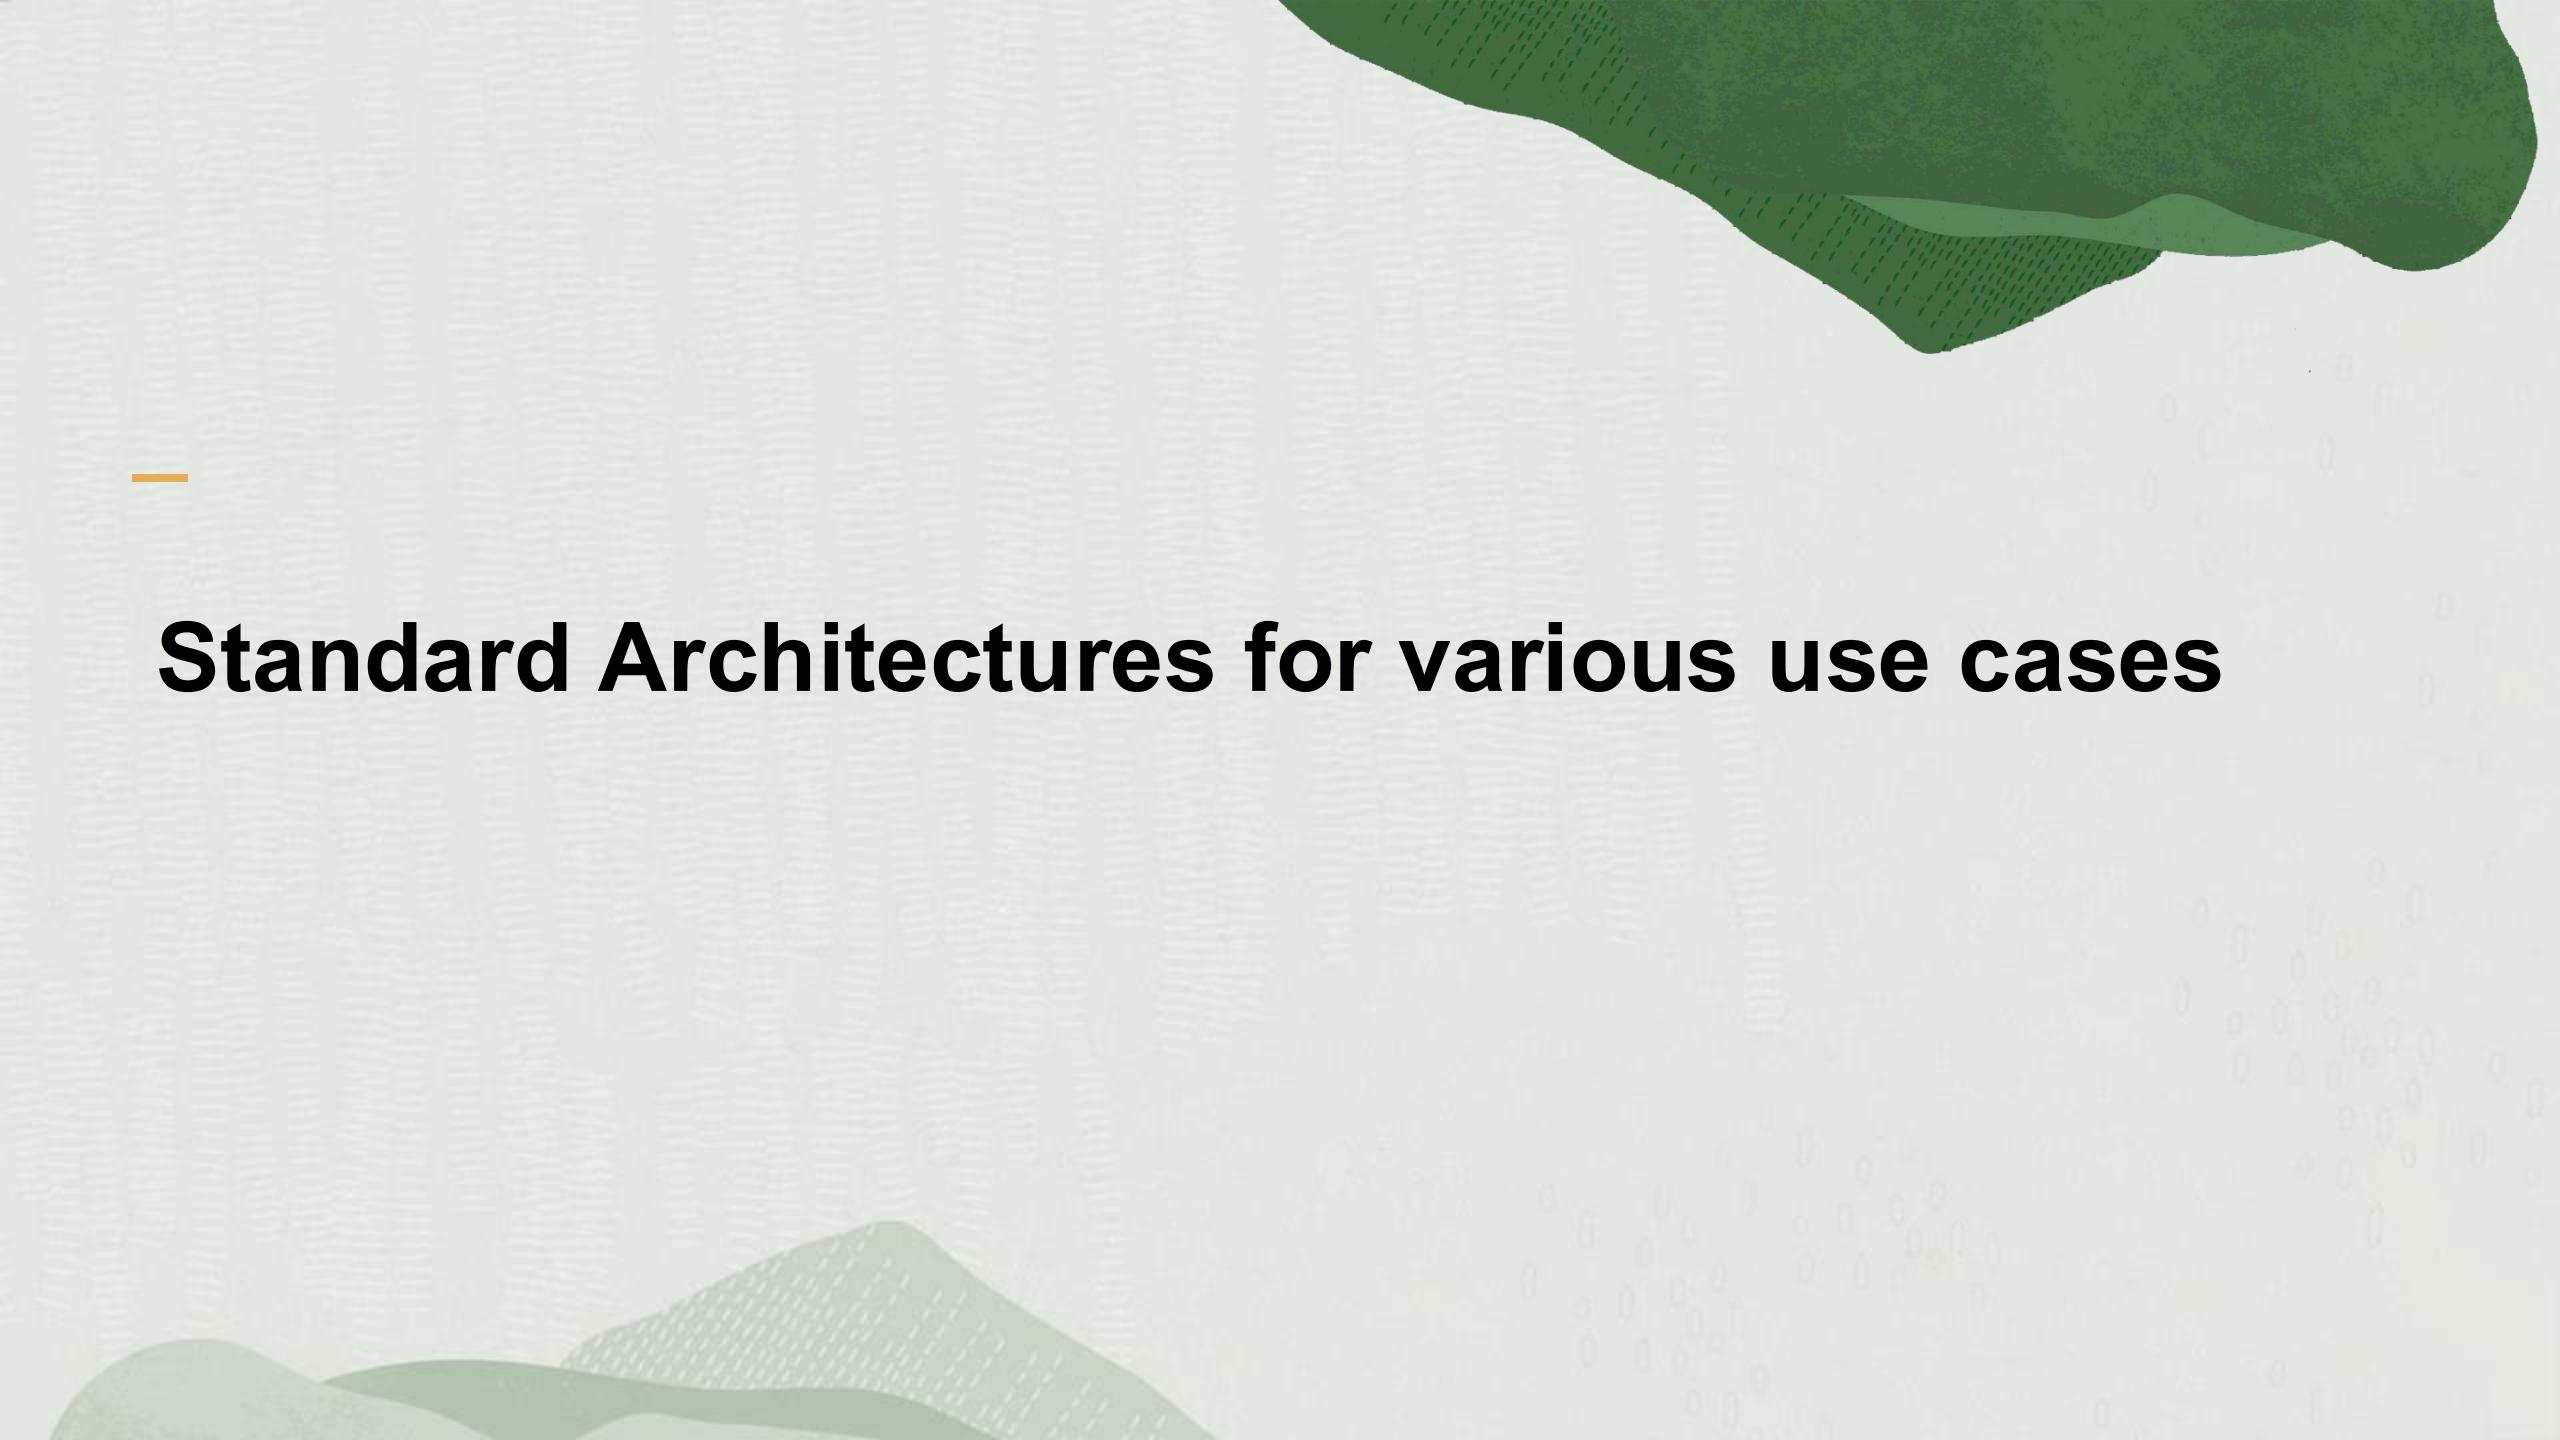
Data and Infrastructure Strategy: Many organizations struggle with siloed, rapidly changing, or low-quality data, which impacts the performance of AI models. It's essential to have a solid data and infrastructure strategy to support the demands of Gen AI.

Regulatory Compliance and Ethical Considerations

Chat Session Management

Guidelines that could be useful

- 1. Ensure User Privacy and Data Security.**
- 2. Maintain Transparency.**
- 3. Promote Ethical Interactions.**
- 4. Manage Expectations:** cut-off knowledge date and its inability to access or retrieve real-time information without specific tools.
- 5. Feedback Mechanism.**



Standard Architectures for various use cases

Generative Adversarial Networks (GANs)

- Excellent for generating photorealistic images, art, and even enhancing low-resolution images.

Use Case

- Training can be unstable and often requires careful tuning of hyperparameters.

Weaknesses

Strengths

- Capable of producing high-quality and highly realistic outputs.

Compared to other models, GANs can generate the most visually compelling results but are often harder to train.

Variational Autoencoders (VAEs)

-
- Use Case
- Used for image generation, but with a focus on encoding an input into a latent space and then reconstructing it.

Strengths

- More stable training than GANs and can learn to represent complex probability distributions.

- Weaknesses
- Tend to produce less sharp images compared to GANs.

VAEs are more stable and easier to train than GANs but do not match the output quality of GANs in terms of image crispness.

Long Short-Term Memory Networks (LSTM)

- Suitable for time-series prediction, music composition, and text generation with a focus on sequences.

Use Case

- Struggles with long-range dependencies and is less effective than transformer models for longer sequences.

Weaknesses

Strengths

- Good at capturing time-dependent properties in sequential data.

LSTMs are more efficient than transformers but lack the same performance on complex tasks requiring understanding of long contexts.

Auto-Regressive Models (e.g., PixelRNN, WaveNet)

- PixelRNN is used for image generation, while WaveNet is used for generating audio, such as human-like speech or music.

Use Case

- The sequential nature of prediction can be slow and computationally expensive

Weaknesses

Strengths

- Can produce high-quality outputs by modelling the probability distribution of a sequence one element at a time.

They are excellent in their respective domains (images and audio) but are generally slower than parallelizable models due to their sequential processing nature.

Comparative Summary



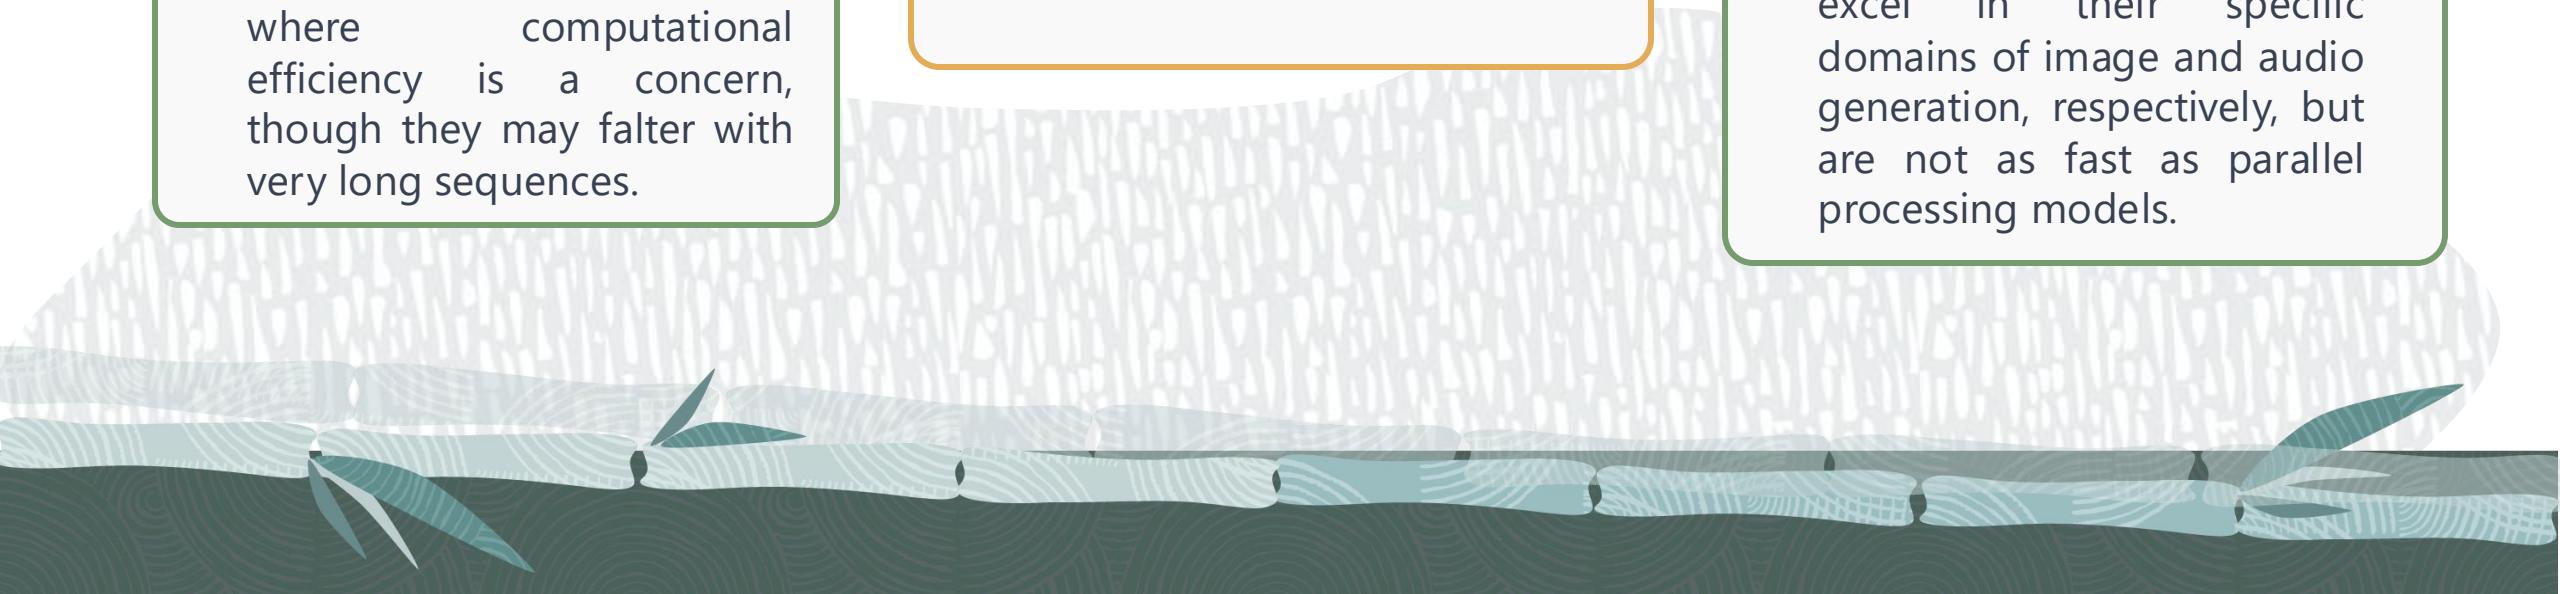
- **GANs** are the best for generating realistic images but are difficult to train.

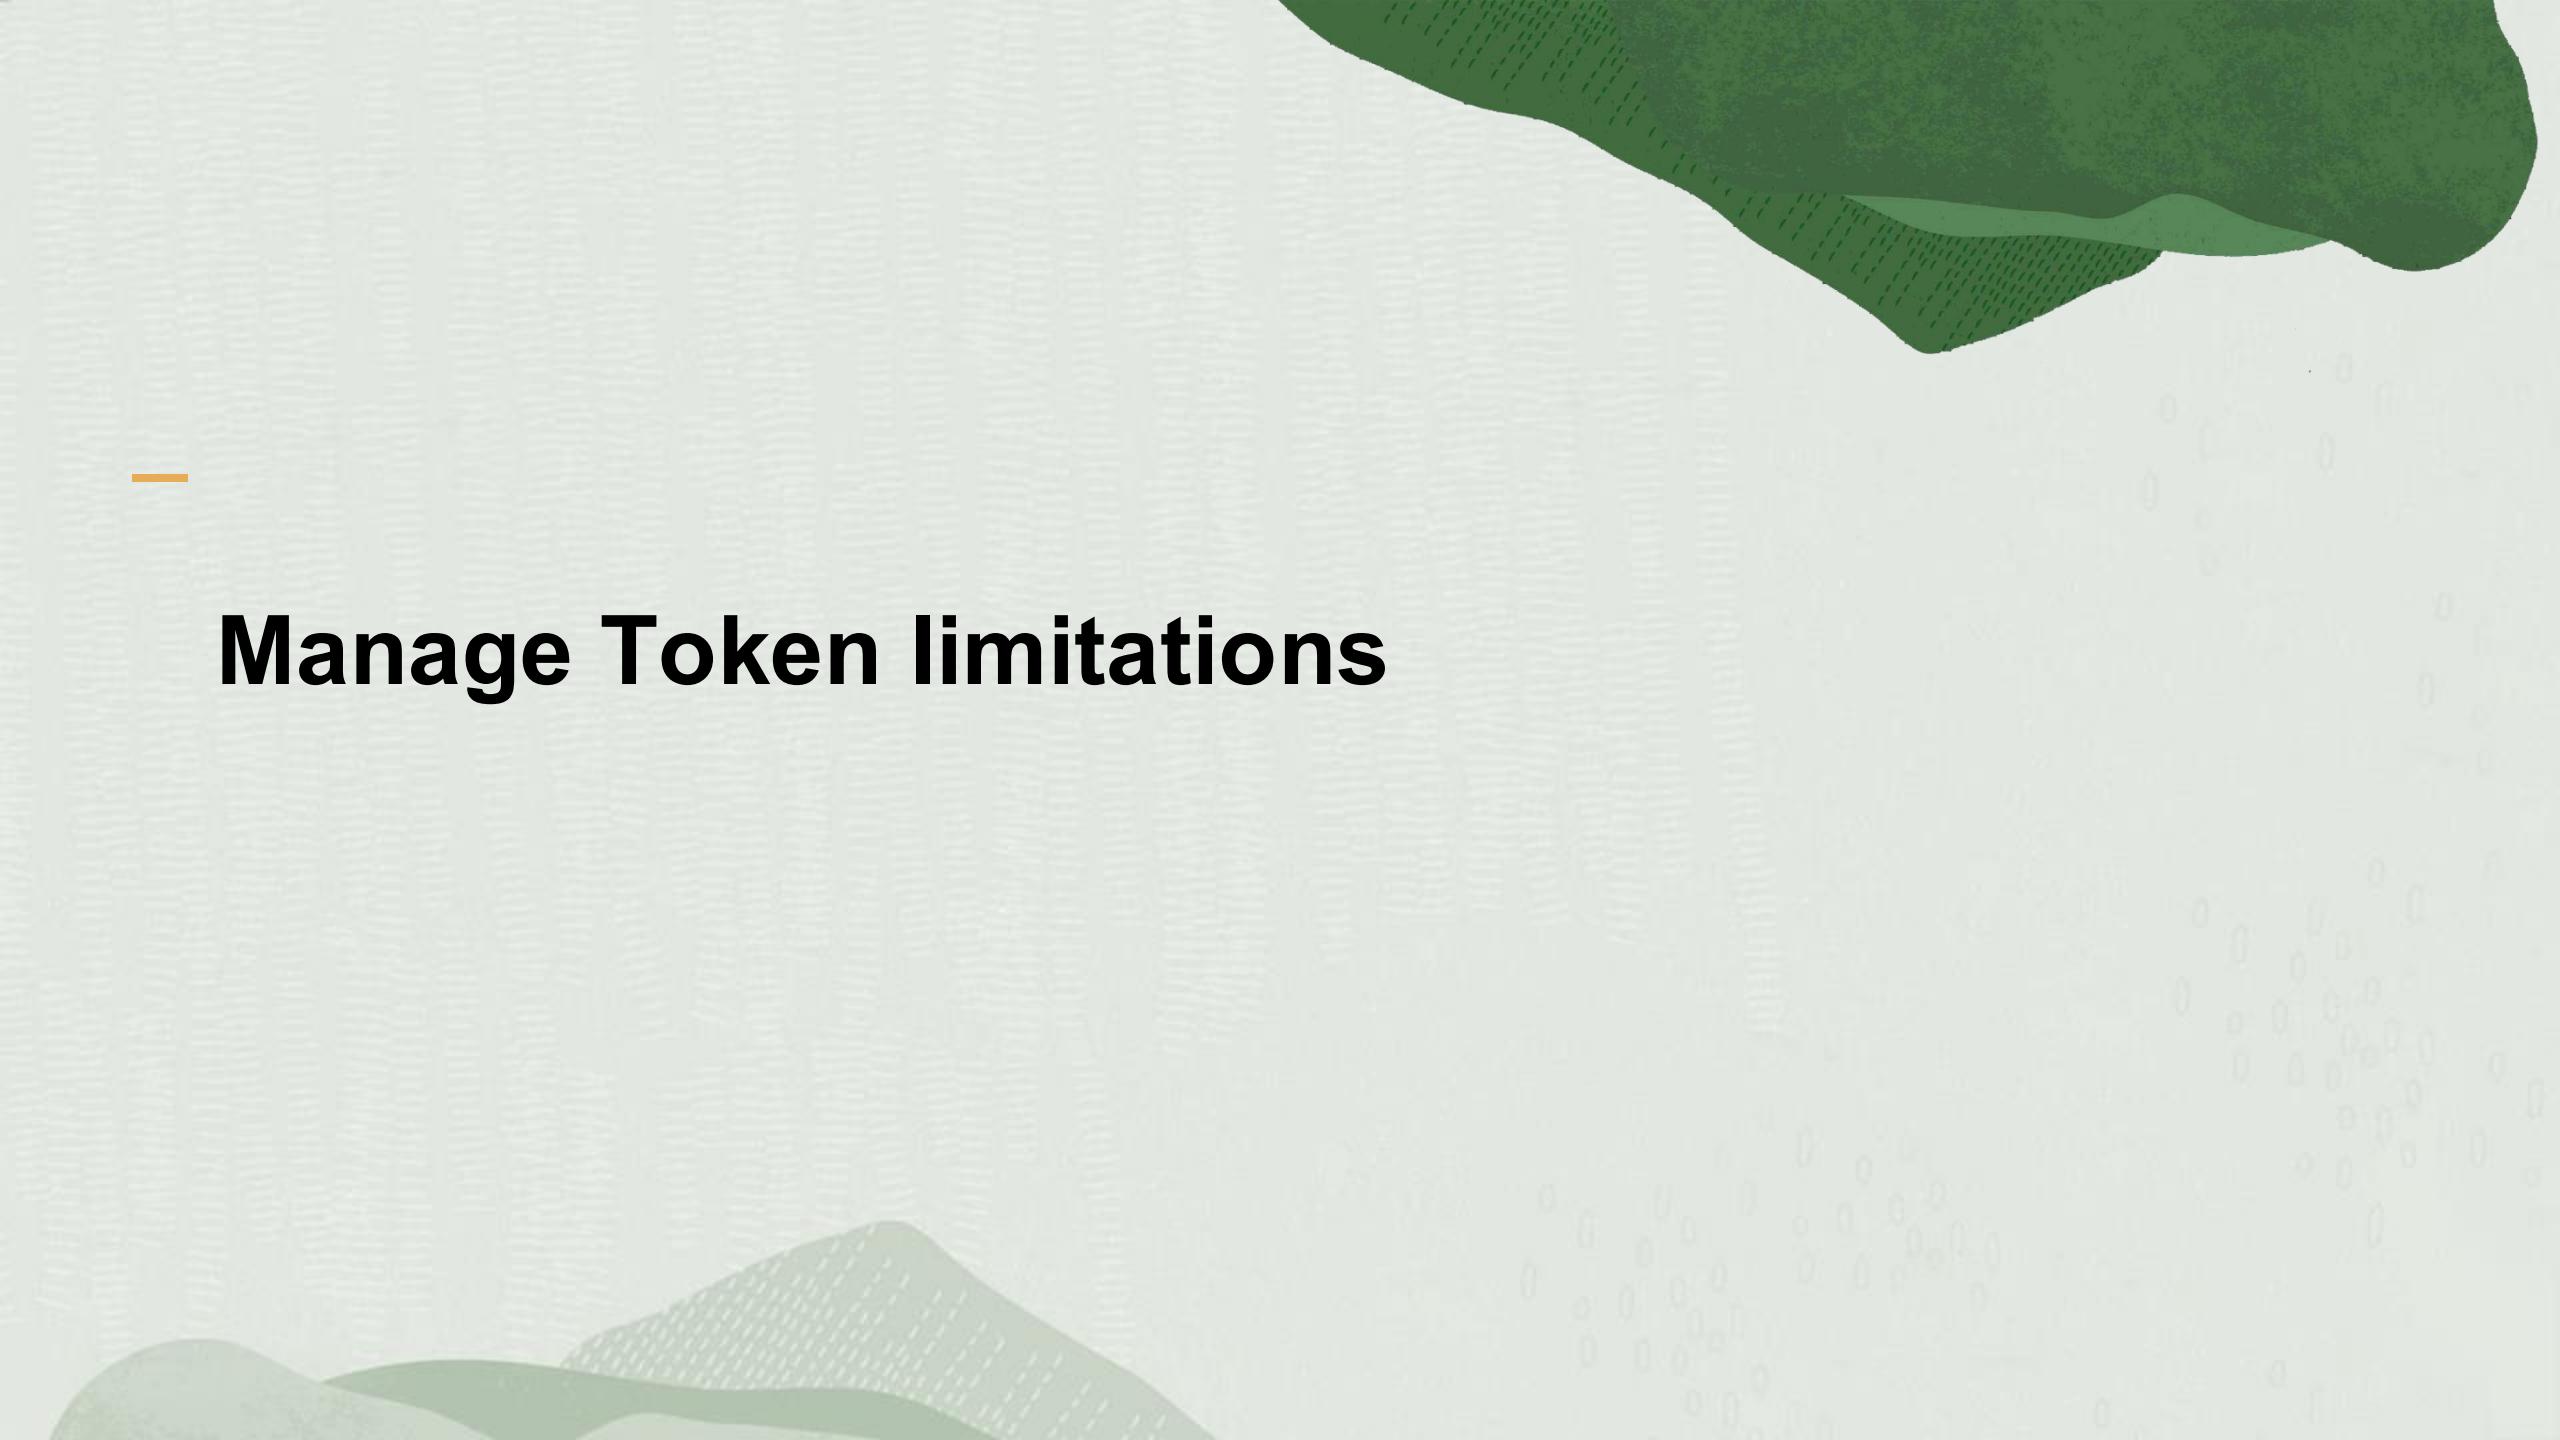
- **LSTMs** are a go-to for simpler sequential tasks where computational efficiency is a concern, though they may falter with very long sequences.

- **VAEs** offer a balance between image quality and training stability and are excellent for tasks that require understanding the underlying data distribution.

- **Transformers** are the gold standard for text-related tasks requiring the understanding of complex contexts but come with high computational costs.

- **Auto-Regressive Models** like PixelRNN and WaveNet excel in their specific domains of image and audio generation, respectively, but are not as fast as parallel processing models.





Manage Token limitations

Manage Token limitations



Summarize Content

Before inputting content into LLM, summarize it to reduce the number of tokens.



Chunking

Process each chunk individually and then synthesize the outputs.

Focused Prompts

Craft your prompts to be as focused as possible.

Iterative Refinement

Iteratively refine the output by asking follow-up questions.

Use Other Models

Consider using smaller models wherever possible, that are more token-efficient.

Increase Efficiency with Commands

Use the model's understanding of commands to execute complex tasks in a token-efficient manner.

Manage Token limitations



Pre-Processing

Use **external tools** for pre-processing tasks like extracting text from images or simplifying sentences to make them more concise.

Post-Processing

After getting the output, use post-processing to stitch together and make sense of the information if you had to break it into parts.

Pipeline Approaches

Create a pipeline that uses different models for different tasks, reserving the large language model for the most complex parts of the task.

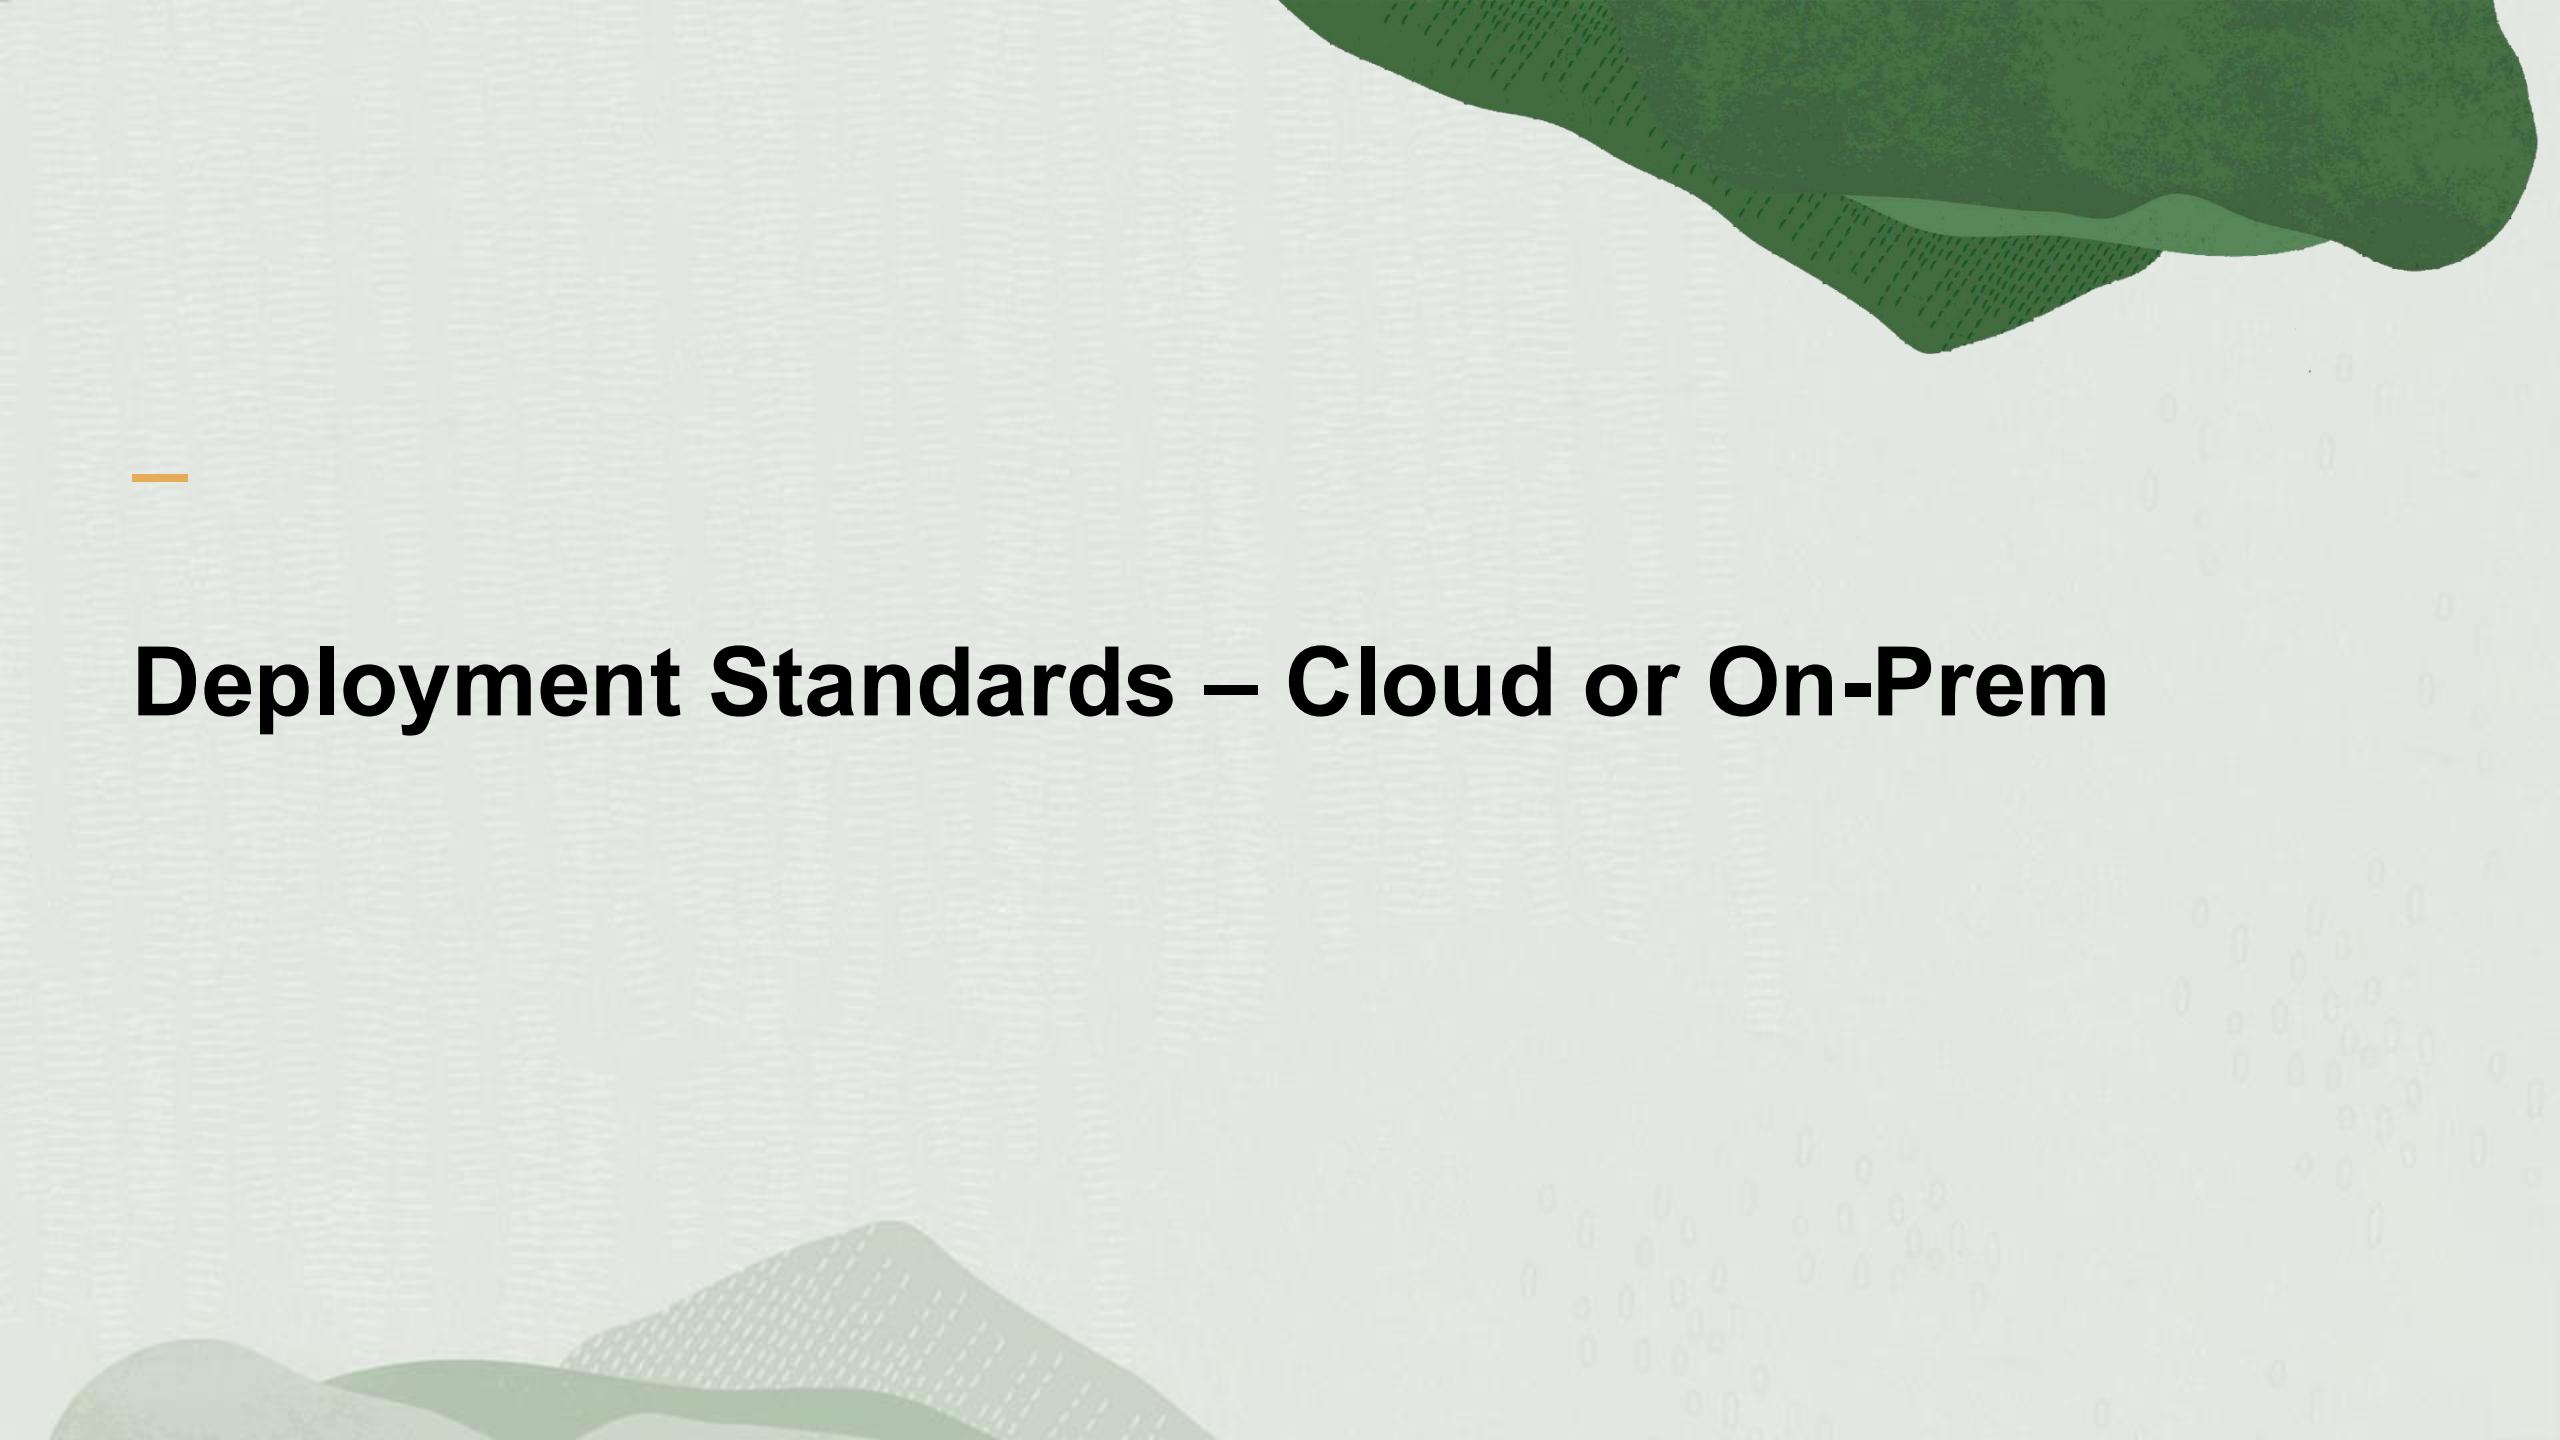
Cache Responses

If you are likely to ask the same or similar questions, cache the responses so you don't have to use tokens for the same query again.

Optimize API Usage

If you are using an API, ensure that your API calls are optimized to send and receive as much relevant information as possible within the token limits.





Deployment Standards – Cloud or On-Prem

Deployment Standards – Cloud or On-Prem

- According to [451 Research](#), 90% of companies use cloud services in some form.
- However, the same report found that 60% of workloads are still run on-premise, indicating a balance between the two.
- In terms of costs, **OCI** can range from a few hundred dollars per month to tens of thousands of dollars per month depending on size of projects.
- However, on-premise solutions can have an upfront cost of several thousand dollars, with additional ongoing costs for maintenance and updates.

Cloud vs. On-Premise Hosting for AI Applications

- These are often compared to renting vs. buying a home.
- Cloud hosting is a lot like renting; the stay of AI applications is as long as the contract terms dictate.
- The maintenance of the hardware is the responsibility of the hosting provider.
- On-premise hosting, on the other hand, is like buying a home; the application can stay on the hardware as long as business requires it.

Cloud vs. On-Premise

Scalability

- On-premise hosting offers complete control over the hardware, which means that the administrators of a company can tightly control updates.
- But on-premise hosting does require advanced planning to scale hardware.
- This is because it requires time to gather the necessary data for updating it.
- Cloud resources can be rapidly adjusted to accommodate specific demands and increase the scalability of hardware.

Cloud vs. On-Premise

Security

- Full control over data stored on enterprise premises.
- Hosting providers must keep their systems updated and data encrypted to avoid breaches.
- Still, your company can't be sure where your data is stored and how often it is backed up.

Cloud vs. On-Premise

Data Gravity

“The ability of data to attract applications, services, and other data towards itself”

It is among the most important factors to be considered while choosing between cloud and on-premise platforms.

- Considering the costs of training neural networks, companies may want to deploy their AI applications on-premises.
- If the data required to build AI applications resides on the cloud, then it's best to deploy applications there.
- The location of the largest source of data for an enterprise determines the location of its most critical applications.
- Oracle cloud offering “**Cloud@Customer**” can prove to be best here.

The Case for Cloud-Based AI Applications

- Instead of building out a massive data centre, you can use the infrastructure someone else already maintains.
- One reason why AI has become so pervasive is cloud providers offering plug-and-play AI cloud services, with enough compute power and pre-trained models to launch AI applications.
- In many cases the pre-trained models or storage requirements of the cloud can be cost-prohibitive; higher GPU counts get expensive fast and training large datasets on the public cloud can be too slow.
- Still, the cloud can often be the best option in terms of “testing the waters” of AI and experimenting with which AI initiatives work best for an organization.

The Case for On-Premise AI

- There's a whole ecosystem of tools built for on-premise infrastructure that can work with mass amounts of compute power—which can be very expensive in the cloud.
- Thus, it is more economical to do this on-premise or prefer a capital expense to an operational expense model.
- If your organization wants to get more involved in this or roll-out AI at scale, then it may make more sense to invest in on-premises infrastructure instead of consuming cloud-based services.

Local Chat

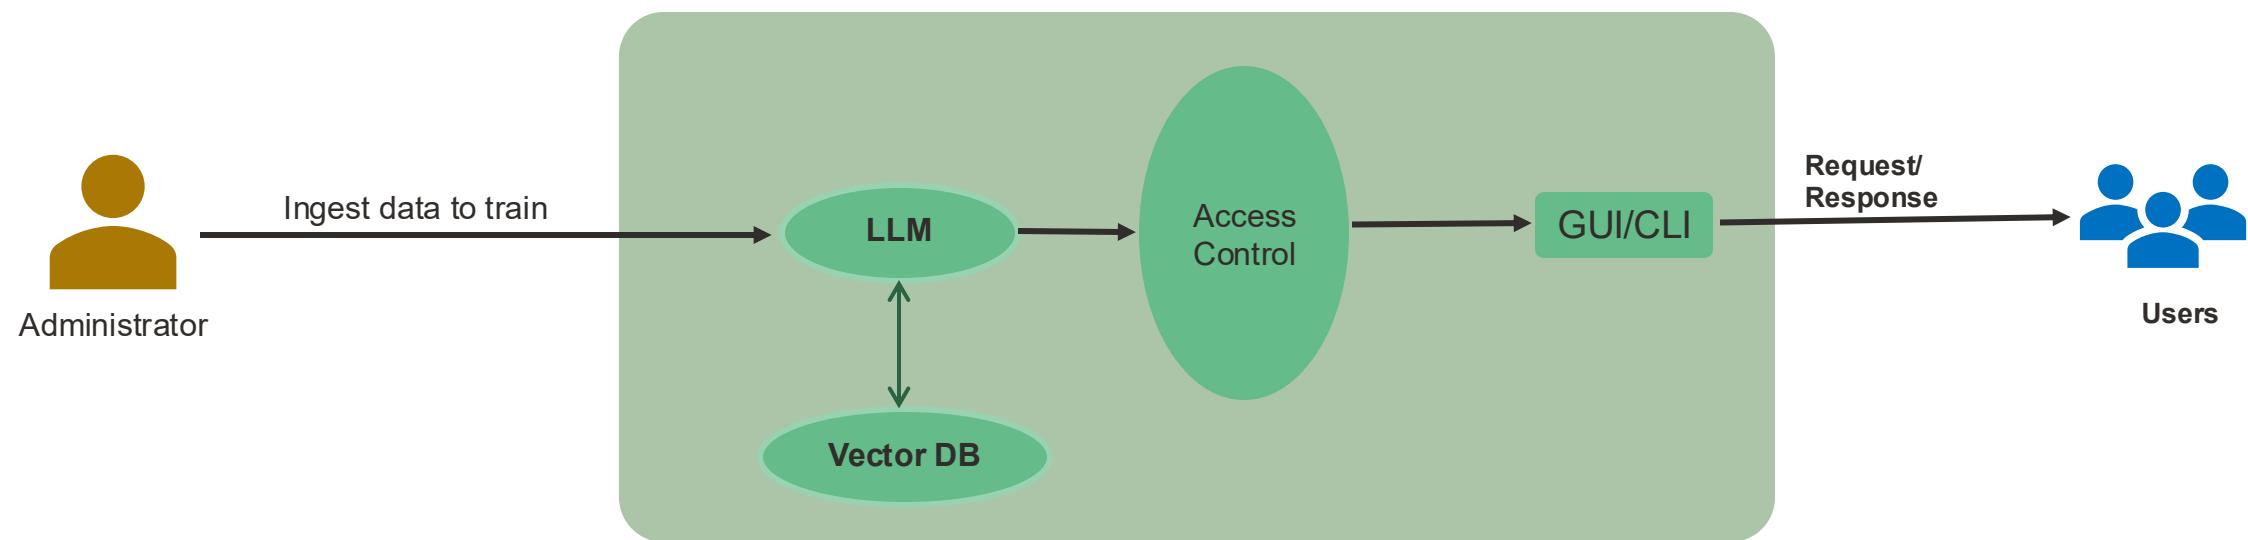
What is LocalChat?

- Local Chat allows you tap into the capabilities of LLMs while prioritizing privacy and security.
- Local Chat operates within an organization's own servers and data centres.
- Use Cases
 - **Knowledge Management.**
 - **Customer Support.**
 - **Content Creation.**
 - **Data Analysis.**
 - **Automate Workflows**

Components of Local Chat and How Does it Work?

Local Chat is requires multiple components to work together to function.

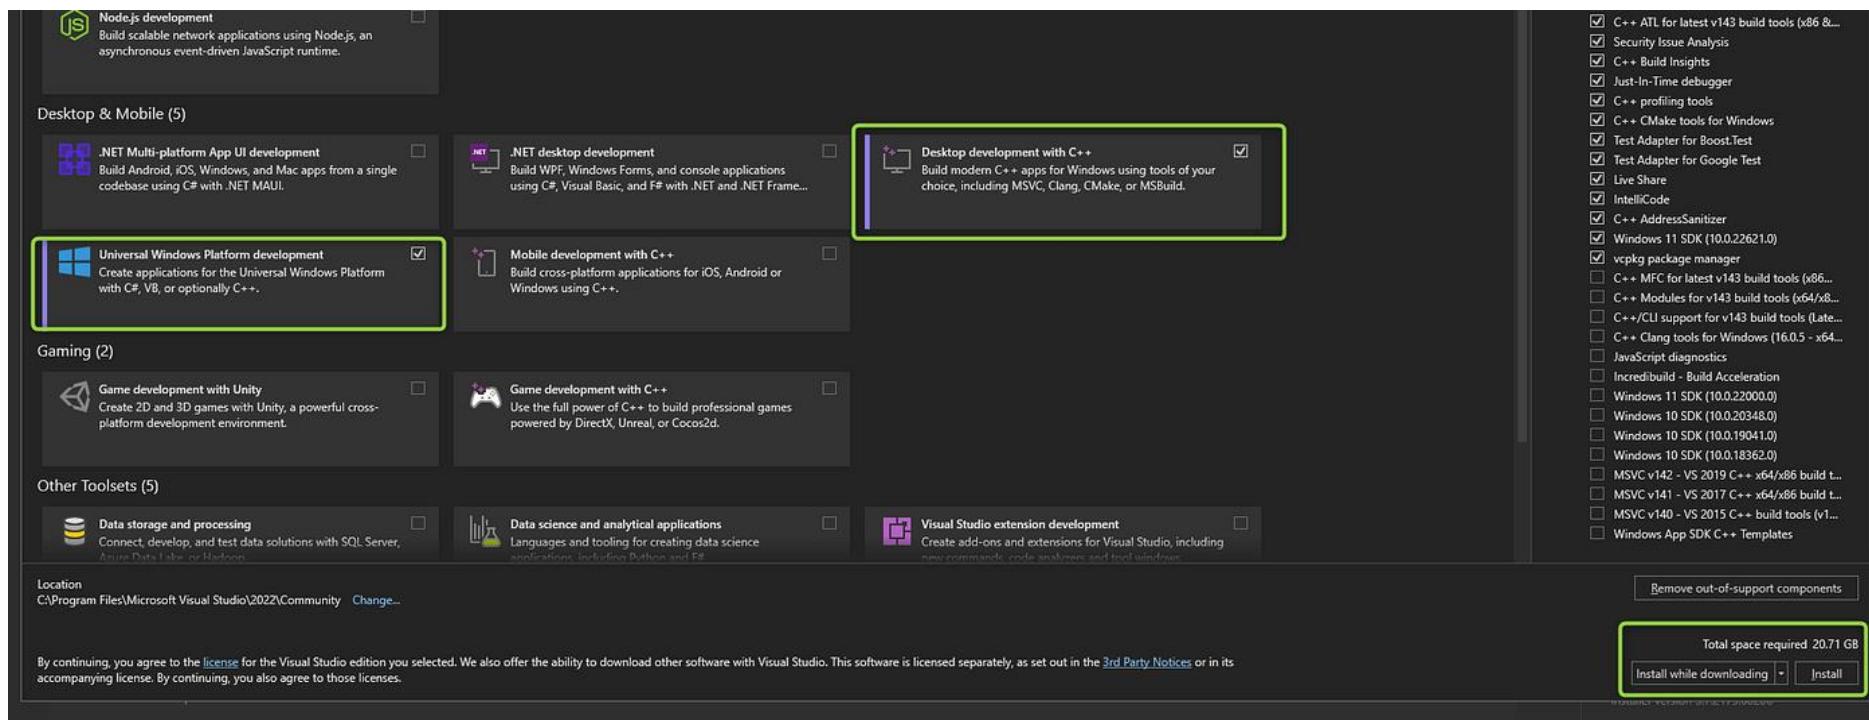
- **Private LLMs:** Local Chat supports proprietary models like [GPT4ALL](#) and [LLAMA](#).



Local Chat - For Windows

Install Visual Studio 2022 Components

1. Universal Windows Platform development
2. Desktop Development with C++

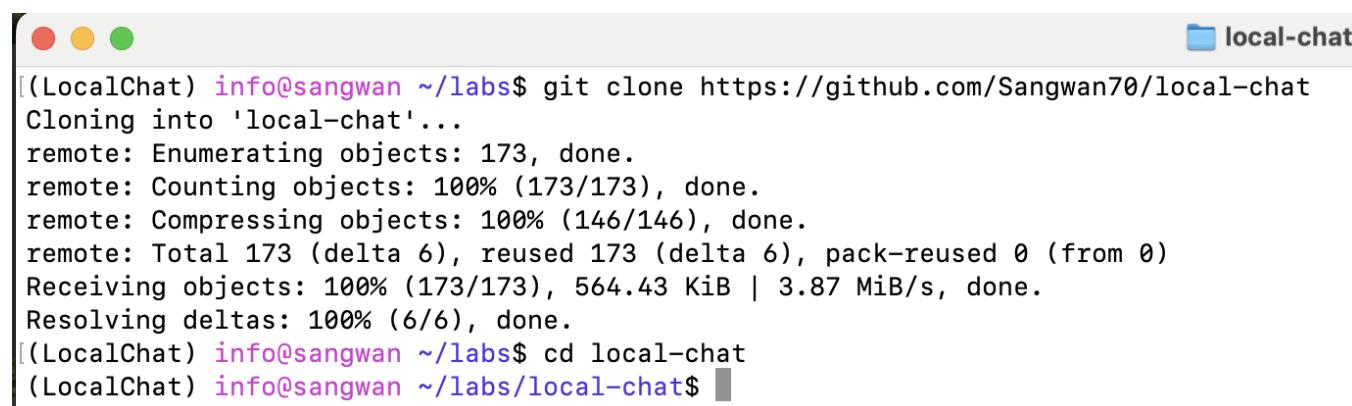


Local Chat – Common for Windows PC/Mac/Linux

Local Chat requires Python 3.11.

- Clone the LocalChat Repository and change to the downloaded directory:

```
git clone https://github.com/Sangwan70/local-chat  
cd local-chat
```



A screenshot of a terminal window on a Mac OS X system. The window title is "local-chat". The terminal shows the command "git clone https://github.com/Sangwan70/local-chat" being run, followed by the output of the cloning process. After the clone is complete, the user runs "cd local-chat" to change into the repository directory. The terminal window has its characteristic red, yellow, and green close buttons at the top left.

```
(LocalChat) info@sangwan ~/labs$ git clone https://github.com/Sangwan70/local-chat  
Cloning into 'local-chat'...  
remote: Enumerating objects: 173, done.  
remote: Counting objects: 100% (173/173), done.  
remote: Compressing objects: 100% (146/146), done.  
remote: Total 173 (delta 6), reused 173 (delta 6), pack-reused 0 (from 0)  
Receiving objects: 100% (173/173), 564.43 KiB | 3.87 MiB/s, done.  
Resolving deltas: 100% (6/6), done.  
(LocalChat) info@sangwan ~/labs$ cd local-chat  
(LocalChat) info@sangwan ~/labs/local-chat$
```

Local Chat for Windows PC/Mac/Linux

Import the LocalChat into an IDE if Needed

- Open terminal and Install Poetry for dependency management.

```
pip install -qU poetry
```

- Install and Run Your Desired Setup

- LocalChat allows customization of the setup, from fully local to cloud-based, by deciding the modules to use.

- **For Windows:**

- *If you are using Windows, you'll need to set the env var, PGPT_PROFILES, in a different way, for example:*

```
# Powershell  
PGPT_PROFILES="ollama"  
make run
```

Local Chat for Windows PC/Mac/Linux

Ollama-Powered Local setup

- Ollama provides local LLM and Embeddings super easy to install and use, abstracting the complexity of GPU support.
- Go to ollama.ai and follow the instructions to install Ollama on your machine.
- Start Ollama service. it will start a local inference server, serving both the LLM and the Embeddings):

ollama serve

```
info@MacBook-Pro private-gpt % pip install -qU poetry
info@MacBook-Pro private-gpt % ollama serve
2024/12/05 22:14:54 routes.go:1197: INFO server config env="map[HTTPS_PROXY: HTTP_
PU_OVERHEAD:0 OLLAMA_HOST:http://127.0.0.1:11434 OLLAMA_KEEP_ALIVE:5m0s OLLAMA_LLM
QUEUE:512 OLLAMA_MODELS:/Users/info/.ollama/models OLLAMA_MULTIUSER_CACHE:false OL
IGINS:[http://localhost https://localhost http://localhost:* https://localhost:* h
http://0.0.0.0 https://0.0.0.0 http://0.0.0.0:* https://0.0.0.0:* app:///* file:///*
tpproxy: https_proxy: no_proxy:]"
time=2024-12-05T22:14:54.539+05:30 level=INFO source=images.go:753 msg="total blob
time=2024-12-05T22:14:54.540+05:30 level=INFO source=images.go:760 msg="total unus
time=2024-12-05T22:14:54.542+05:30 level=INFO source=routes.go:1248 msg="Listening
time=2024-12-05T22:14:54.542+05:30 level=INFO source=common.go:135 msg="extracting
ama3734876440/runners
time=2024-12-05T22:14:54.595+05:30 level=INFO source=common.go:49 msg="Dynamic LLM
time=2024-12-05T22:14:54.616+05:30 level=INFO source=types.go:123 msg="inference c
="21.3 GiB" available="21.3 GiB"
```

LocalChat for Windows PC/Mac/Linux

Install the models to be used

- The default `settings-ollama.yaml` is configured to user llama 3.2 LLM and nomic-embed-text Embeddings (~275MB)
- By default, PGPT will automatically pull models as needed.
- In any case, if you want to manually pull models, run:

```
ollama pull llama3.1
```

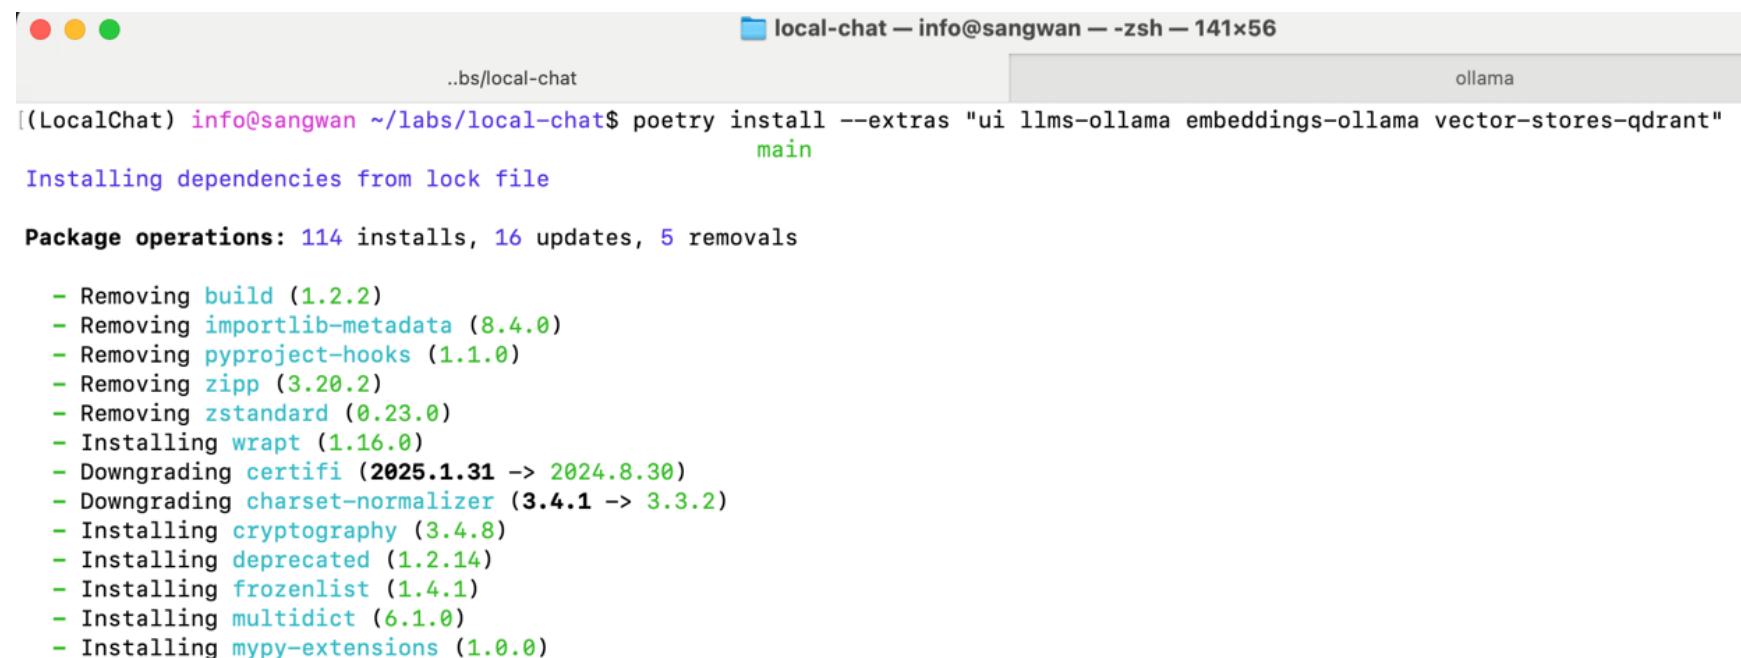
```
ollama pull nomic-embed-text
```

Local Chat for Windows PC/Mac/Linux

Install LocalChat

- Once done, on a different terminal, install LocalChat:

```
poetry install --extras "ui llms-ollama embeddings-ollama vector-stores-qdrant"
```



A screenshot of a terminal window titled "local-chat — info@sangwan — -zsh — 141x56". The window shows the command "poetry install --extras "ui llms-ollama embeddings-ollama vector-stores-qdrant"" being run. The terminal output indicates it is installing dependencies from a lock file, with 114 installs, 16 updates, and 5 removals. A list of packages being removed and installed is shown.

```
..bs/local-chat
(localChat) info@sangwan ~/labs/local-chat$ poetry install --extras "ui llms-ollama embeddings-ollama vector-stores-qdrant"
main
Installing dependencies from lock file

Package operations: 114 installs, 16 updates, 5 removals

- Removing build (1.2.2)
- Removing importlib-metadata (8.4.0)
- Removing pyproject-hooks (1.1.0)
- Removing zipp (3.20.2)
- Removing zstandard (0.23.0)
- Installing wrapt (1.16.0)
- Downgrading certifi (2025.1.31 -> 2024.8.30)
- Downgrading charset-normalizer (3.4.1 -> 3.3.2)
- Installing cryptography (3.4.8)
- Installing deprecated (1.2.14)
- Installing frozenlist (1.4.1)
- Installing multidict (6.1.0)
- Installing mypy-extensions (1.0.0)
```

Local Chat for Windows PC/Mac/Linux

Run LocalChat

- Make sure you have a working Ollama running locally before running the following command.

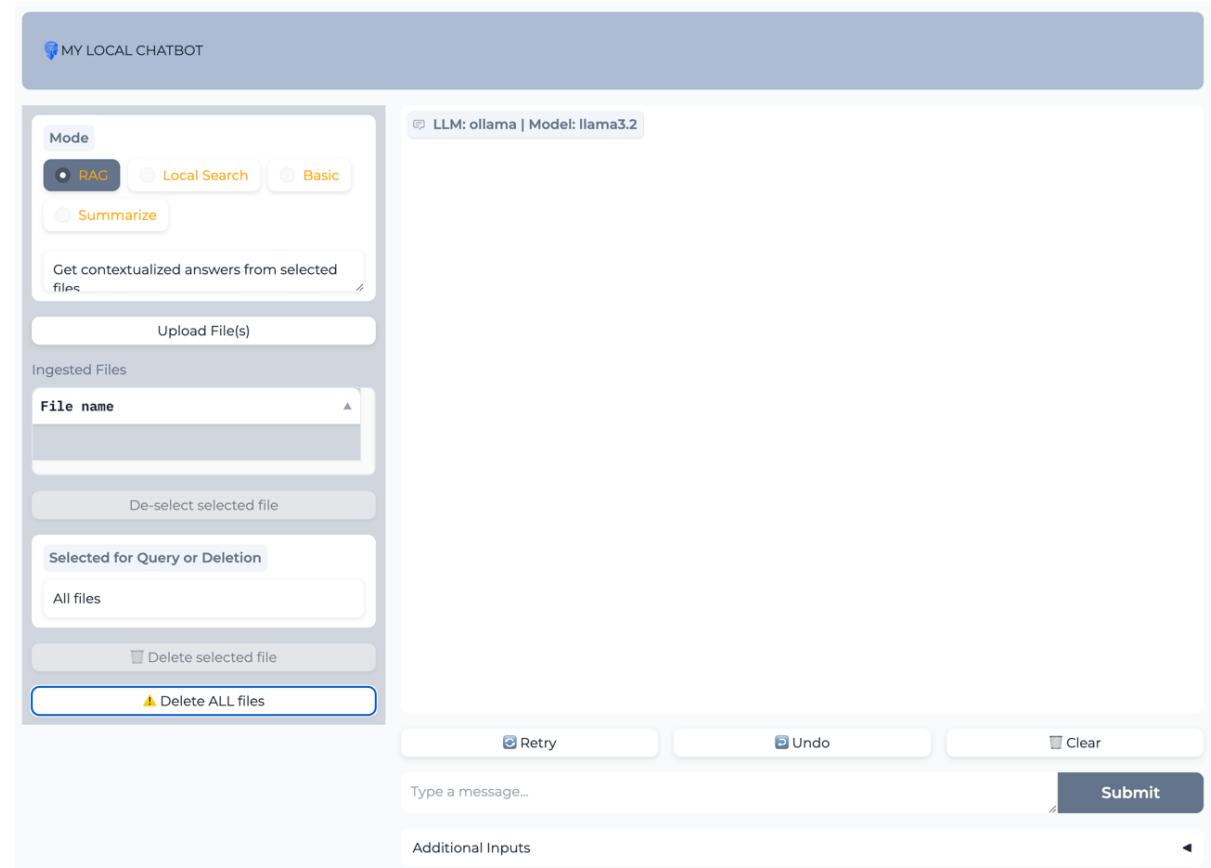
```
PGPT_PROFILES=ollama make run
```

```
22:19:05.400 [INFO    ] private_gpt.components.llm.llm_component - Initializing the LLM in mode=ollama
22:19:05.486 [INFO    ] httpx - HTTP Request: GET http://localhost:11434/api/tags "HTTP/1.1 200 OK"
22:19:05.488 [INFO    ] httpx - HTTP Request: GET http://localhost:11434/api/tags "HTTP/1.1 200 OK"
22:19:06.636 [INFO    ] private_gpt.components.embedding.embedding_component - Initializing the embedding model in mode=ollam
22:19:06.677 [INFO    ] httpx - HTTP Request: GET http://localhost:11434/api/tags "HTTP/1.1 200 OK"
22:19:06.679 [INFO    ] httpx - HTTP Request: GET http://localhost:11434/api/tags "HTTP/1.1 200 OK"
22:19:06.679 [INFO    ] llama_index.core.indices.loading - Loading all indices.
22:19:06.679 [INFO    ] private_gpt.components.ingest.ingest_component - Creating a new vector store index
Parsing nodes: 0it [00:00, ?it/s]
Generating embeddings: 0it [00:00, ?it/s]
22:19:16.377 [INFO    ] private_gpt.ui.ui - Mounting the gradio UI, at path=/
22:19:16.828 [INFO    ] uvicorn.error - Started server process [6180]
22:19:16.828 [INFO    ] uvicorn.error - Waiting for application startup.
22:19:16.828 [INFO    ] uvicorn.error - Application startup complete.
22:19:16.829 [INFO    ] uvicorn.error - Uvicorn running on http://0.0.0.0:8001 (Press CTRL+C to quit)
```

Local Chat for Windows PC/Mac/Linux

Access LocalChat

- LocalChat will use the already existing settings-ollama.yaml, which is already configured to use Ollama LLM and Embeddings, and Qdrant.
- Review it and adapt it to your needs.
- The UI will be available at <http://localhost:8001>



Mode

RAG Local Search Basic

Summarize

Chat with the LLM using its training data. Files are ignored.

Upload File(s)

Ingested Files

File name

responsible-ai-oracle.pdf

De-select selected file

Selected for Query or Deletion

All files

Delete selected file

⚠ Delete ALL files

MY LOCAL CHATBOT

Mode

RAG

Local Search

Basic

Summarize

Get contextualized answers from selected files.

Upload File(s)

Ingested Files

File name

responsible-ai-oracle.pdf

De-select selected file

Selected for Query or Deletion

All files

Delete selected file

⚠ Delete ALL files

LLM: ollama | Model: llama3.2

Tell me Something about Oracle's Responsible AI Policy.

The text mentions "Oracle's Guide to Ethical Considerations in AI Development and Deployment" as a resource for responsible AI practices, but it does not provide specific details about their policy itself. However, I can tell you that the guide is intended to help organizations develop and deploy AI systems in an ethical manner, with considerations such as avoiding bias, ensuring transparency, and promoting accountability.

Sources:

1. responsible-ai-oracle.pdf (page 8)
2. responsible-ai-oracle.pdf (page 9)



Retry

Undo

Clear

Type a message...

Submit

Additional Inputs



Thank You